

# SyRust: Automatic Testing of Rust Libraries with Semantic-Aware Program Synthesis

Anonymous Author(s)

## Abstract

Rust’s type system ensures the safety of Rust programs; however, programmers can side-step some of the strict typing rules by using the `unsafe` keyword. A common use of `unsafe` Rust is by libraries. Bugs in these libraries undermine the safety of the entire Rust program. Therefore, it is crucial to thoroughly test library APIs to rule out bugs. Unfortunately, such testing relies on programmers to manually construct test cases, which is an inefficient and ineffective process.

The goal of this paper is to develop a methodology for automatically generating Rust programs to effectively test Rust library APIs. The main challenge is to synthesize *well-typed* Rust programs to account for proper chaining of API calls and Rust’s ownership type system and polymorphic types. We develop a program synthesis technique for Rust library API testing, which relies on a novel logical encoding of typing constraints from Rust’s ownership type system. We implement SyRust, a testing framework for Rust libraries that automatically synthesizes semantically valid test cases. Our experiments on 30 popular open-source Rust libraries found 4 new bugs.

## 1 Introduction

Despite being only a decade old, the Rust language has found high popularity in a wide range of domains such as operating systems, embedded devices, and high-performance web frameworks [1]. A key ingredient of Rust’s success is the ownership and variable lifetime system, which guarantees any program passing the type check is free of unsafe behaviors such as use-after-free, double-free, and data races. To recover some expressiveness, Rust also allows programmers to disable some of the safety checks for code segments labeled with the `unsafe` keyword. A common use of `unsafe` Rust is by libraries, which encapsulate `unsafe` code and provide abstractions to safe Rust. Allowing `unsafe` code has inevitably led to vulnerabilities in Rust applications. The Rust Standard Library was found to be vulnerable in 2018 [2] and recent surveys [24, 31] further demonstrate the danger of using `unsafe` Rust.

Effort has been made to help programmers mitigate the risks stemmed from `unsafe` code, for instance, coding guidelines [5] and the Stacked Borrows project that can determine if a Rust program exhibits buggy behavior [19]. Fundamentally, it is necessary for programmers to thoroughly test their

libraries that include `unsafe` code. Unfortunately, such testing relies on programmers to manually construct test cases, which can be inefficient and ineffective.

Our goal is to automatically generate test cases for Rust libraries. The main challenges stem from Rust type system features such as polymorphism and variable ownership and lifetime restrictions that make Rust safe. The Rust compiler will reject a large portion of programs that are generated using traditional random search techniques like Randoop [23] and RESTler [7]. Instead, we take inspiration from the recent success of constraint solving based program synthesis techniques like H+ [14] and SyPet [13]. These techniques take in a set of API specifications and synthesize valid programs by solving a constraint formula that encodes the set of all bounded-length programs that can be built using the given APIs. In this paper, we propose novel, Rust-specific constraints that encode the above mentioned Rust type system features, allowing us to synthesize valid Rust programs.

Concretely, we develop a semantic-aware synthesis algorithm that encodes typing constraints including ownership and lifetime constraints by keeping track of the typing contexts and which variables are active and thus can be used as arguments to APIs at each program point. Programs that satisfy these constraints are most likely to pass Rust type checker. Further, this algorithm facilitates the generation of valid API call chains, which are necessary for exposing subtle bugs that would have been missed by testing one API at a time. To make use of APIs with polymorphic types effectively, we propose a hybrid technique to effectively concretize polymorphic object constructors while avoiding combinatorial explosions through early pruning. Furthermore, we develop a type refinement algorithm that leverages error reports provided by the Rust compiler. In the end, we achieve a synthesis algorithm that can effectively generate programs that make use of APIs with polymorphic types without type errors.

We implement the proposed techniques in SyRust, a scalable, synthesis driven, testing framework for Rust libraries. We evaluated SyRust on 30 of popular open source Rust libraries, and found 4 bugs, including a double-free that requires a complex sequence of five calls to trigger. Furthermore, our evaluation shows that SyRust is highly effective at generating test programs while minimizing compiler errors. Finally, we demonstrate that the semantic awareness and type refinement is critical to the success of SyRust by observing a significant increase in the number of compiler-rejected test cases when the respective features are turned off. We summarize our contributions as follows.

1. We propose a novel semantic-aware synthesis algorithm that generates valid Rust programs.
2. We improve the state-of-the-art type refinement synthesis technique to accommodate Rust-specific features by leveraging Rust compiler error messages.
3. We implement our synthesis algorithm in SyRust, an automatic testing framework for Rust libraries.
4. We evaluate SyRust on popular Rust libraries and demonstrate that it can generate semantically valid test cases that expose bugs.

## 2 Background and Motivation

We discuss key features of the Rust type system and highlight challenges in automatic generation of valid Rust programs. Considering a simple test case shown in Figure 1 for `Vec<T>`, a vector data structure from the Rust standard library. The type context for each line is shown in the comments. The test function takes as arguments a `String` object (`s:&String`) and a vector of `Strings` (`v :Vec<String>`). Line 3 casts `v` into mutable vector `vm` and line 5 creates a mutable reference `vr` that points to `vm`. Using this reference, `s` is pushed to the vector on line 8. Line 10 calls `into_raw_parts`, a function that destroys the vector and returns a tuple of size 3.

```

1  fn test(s :String, v :Vec<String>) {
2  //context:{s:String,v:Vec<String>}
3  let mut vm = v;
4  //context:{s:String,vm:mut Vec<String>}
5  let vr :&mut Vec<String> = &mut v;
6  //context:{s:String, vm:Vec<String>,
7  //          vr:&mut Vec<String>}
8  vr.push(s);
9  //context:{vm:mut Vec<String>,vr:&mut Vec<String>}
10 let _ = vm.into_raw_parts();
11 //context: {}
12 }
```

Figure 1. An example Rust program

This small well-typed program chains several API calls together and is very similar to a test case generated by SyRust that exposed a bug in `bitvec`, a bitvector library for Rust. (Figure 6). For the rest of this section, we will use this program as an example to review Rust type system and illustrate challenges in automatically generating such a test program.

**Basic Typing and Subtyping Constraints** Rust is statically typed: all variables are assigned a type at compile time and all uses of variables are consistent with their types. For example, if an API takes a string (`String`) as its argument, then only variables (expressions) that have the type `String` can be given as argument to the API. Rust also allows subtyping. For example, a mutable reference to a string `&mut Vec<T>` may be used in place of immutable reference `&Vec<T>`, but not the other way around. This means that our test case generating algorithm needs to keep track of typing contexts and reason about types and subtyping.

**Polymorphism** The type variable `T` in `Vec<T>` can be instantiated with concrete types. To generate test cases for APIs with polymorphic types, our algorithm needs to know how the type variable is instantiated and handle subsequent typing constraints. For example, the push operation seen above has type `(&mut Vec<T>, T)→()`. Our algorithm need to match concrete vector reference `&mut Vec<String>` with the polymorphic input type `&mut Vec<T>` and then understand that the other input type should match the type previously used for instantiating `T`. Naive enumeration of concrete types to instantiate type variables is inefficient; polymorphic types can be instantiated with other polymorphic types (`Vec<T>`, `Vec<Vec<T>>`, `Vec<Vec<Vec<T>>>` and so on), leading to possible infinite instantiation loops. Type variables that are required to be instantiated with types implementing a trait, indicating supports for sets of methods, is yet another constraint we need to consider.

**Variable Lifetime and Ownership** A lifetime of a variable is a region in the code in which this variable is in scope. If a variable is used, then the compiler ensures that this usage occurs within the lifetime of the variable. While the notion of lifetime is common across all imperative languages, Rust employs one that allows it to prevent memory errors. When a variable of non-primitive type is used in an expression, Rust forcefully terminates the lifetime of that variable at that line as it moves the value out of the variable and into the input variable of the function. This maintains an invariant that only *one* variable token owns a given memory location. Guaranteeing single ownership makes reasoning about when an object can be deallocated easy. For example, the lifetime of `s` is terminated at line 8; if we were to call `vr.push(s)`; again on line 9, the program will no longer type check. Consequently, our algorithm needs to explicitly handle lifetime and ownership constraints.

**Relationship Between Lifetimes** Sometimes, a variable's lifetime depends on another variable. For example `vr` is a reference to `vm`. To avoid creating a potential use-after-free bug, the lifetime of `vr` must be strictly contained in the lifetime of `vm`. If we swap the last 2 lines of Figure 1 (below) this program will not type check because when `vm` is destroyed by `into_raw_parts`, `vr` is also removed from the type context.

```

1  //context:{s:String, vm:Vec<String>,
2  //          vr:&mut Vec<String>}
3  let _ = vm.into_raw_parts();
4  //context:{s:String}
5  vr.push(s); // vr not found
```

**Borrowing and References** To prevent data races, Rust enforces the rule that only one mutable reference can be active for any location in memory. The compiler will reject any attempts to borrow a mutable reference to a memory location while another mutable reference is active. For example, the following program attempts to borrow a second mutable reference `vr2`. This does not pass the Rust compiler.

```

1 //context{s:String,v:Vec<String>}
2 let mut vm = v;
3 //context:{s:String,vm:mut Vec<String>}
4 let vr :&mut Vec<String> = &mut vm;
5 //context:{s:String, vm:Vec<String>,
6 //      vr:&mut Vec<String>}
7 let vr2 :&mut Vec<String> = &mut vm;
8 //context: can't have both vr and vr2
9 vr.push(s);

```

Even if `vr2` is an immutable reference, the program still causes a type error; because mutable and immutable references for one memory location cannot co-exist. A memory location can have many immutable only references. Our algorithm need to encode these constraints on references.

### 3 SyRust Overview

The architecture of SyRust is shown in Figure 2. SyRust implements an iterative approach and consists of three main components: a semantic-aware test case synthesis engine; an API specification refinement engine, and a test executor.

**Synthesis engine** The synthesis engine takes as inputs a code template and library API typing specifications and returns a set of test cases. The engine implements a novel semantic-aware synthesis algorithm that takes into consideration constraints such as those discussed in Section 2. The algorithm is explained in Section 4.

The API typing specifications are collected from the library to be tested. The code template is manually generated, one for each target the analyst intends to test. An example code template for the vector library is shown in Figure 3. It serves the following purposes. First, it indicates to the synthesizer where to insert the synthesized code via the (`// INSERT`) comments. Second, the main function provides inputs to the test function. Typically, these inputs are assigned types so that they can be used as arguments to the APIs to be tested.

```

1 fn test(s :&String, v :Vec<String>) {
2 //INSERT
3 }
4
5 fn main() {
6 let s_0 : std::string::String = "".to_string();
7 let v_0 : Vec<String> = Vec::new();
8
9 test(&s_0, v_0);
10 }

```

Figure 3. An example template for `Vec<T>`

**Polymorphic API Refinement** Many Rust library APIs and data structures use polymorphic types. To effectively test them, we implement a novel *hybrid* type variable instantiation scheme that eagerly instantiates a small subset of APIs at the beginning and uses compiler errors to lazily refine the remaining API specifications. If a compiler error is encountered when testing generated programs, the polymorphic type refinement module alters the API specification, so

#### Algorithm 1: High-Level Algorithm of SyRust

**Input** : API type signatures  $\mathcal{A}$ , code template  $\mathcal{T}$ , maximum number of lines of code  $m$

**Output**: database of programs and results  $\mathcal{DB}$

1 **Procedure** *SyRust*( $\mathcal{A}, \mathcal{T}, m$ ):

```

2   for  $l \in 1, \dots, m$  do
3      $\varphi \leftarrow \text{genConstraints}(\mathcal{A}, \mathcal{T}, l)$ 
4     while  $\text{SATSolver}(\varphi) \neq \text{UNSAT}$  do
5        $\sigma \leftarrow \text{getModel}(\text{SATSolver}(\varphi))$ 
6        $\mathcal{P} \leftarrow \text{codeGen}(\sigma, \mathcal{T})$ 
7        $\mathcal{R} \leftarrow \text{test}(\mathcal{P})$ 
8       if  $\text{isCompilerError}(\mathcal{R})$  then
9          $\mathcal{A} \leftarrow \text{refine}(\mathcal{A}, \mathcal{R})$ 
10         $\varphi \leftarrow \text{update}(\varphi, \mathcal{A})$ 
11      end
12       $\mathcal{DB} \leftarrow \mathcal{DB} \cup \mathcal{R}$ 
13       $\varphi \leftarrow \varphi \wedge \neg \sigma$ 
14    end
15  end
16  return  $\mathcal{DB}$ 

```

that the same error does not occur again. Based on the error, it first identifies the (polymorphic) API that causes it, then, a new API specification with the instantiated type is added to the API specifications to be used in the next round of synthesis. For example, `Vec::pop` returns `Option<T>`. Initially, we only have the polymorphic version. If we were to use this API with a vector of type `Vec<String>`, then we add a new API with `T` instantiated as `String` for both the inputs and outputs. APIs deemed unfixable will be prevented from being used by the synthesizer. Details of this refinement process will be presented in Section 5.

**Test Executor** The test executor has two roles. First, it compiles the test cases and reports compiler errors. Second, it executes compiled test cases and reports unbounded behavior (bugs). For our implementation, we leverage Miri [22], an interpreter for Rust that flags unbounded behavior (e.g., dereferencing freed memory). Any tool that can satisfy the above two roles can be used.

## 4 Semantics-Aware Synthesis

In this section, we show how we can synthesize code that follows the Rust semantics described in Section 2.

### 4.1 Synthesis Algorithm

Algorithm 1 shows an overview of SyRust's synthesis algorithm. SyRust takes as input a set of API type signatures ( $\mathcal{A}$ ), a code template ( $\mathcal{T}$ ), and a maximum number of lines of code to synthesize ( $m$ ). The goal of this procedure is to synthesize a collection of test programs that follow the Rust semantics and can be successfully compiled by the Rust compiler. To

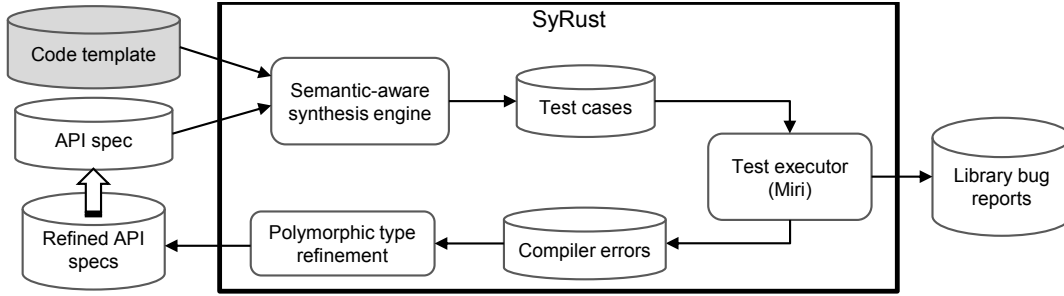


Figure 2. Overview of SyRust

achieve this goal, we synthesize programs of increasing size by encoding the search space of all possible programs of size  $m$  that satisfy Rust's semantic properties as a SAT formula  $\varphi$  (line 3). Each model  $\sigma$  of  $\varphi$  can be translated to a concrete program  $\mathcal{P}$  (line 6) that can be compiled and executed by the test executor (line 7). If a compiler error is seen (line 8), then we analyze the error and detect which APIs were the cause of the compilation error in order to refine  $\mathcal{A}$  such that this error does not occur in the future. This can occur for instance when an API is polymorphic and we did not instantiate it. Regardless of the result, we save the solution and resulting messages from test execution in a database  $\mathcal{DB}$  (line 12). Finally, we block  $\sigma$  to avoid repetition (line 13) and repeat the process until all models are found and the formula becomes unsatisfiable.

This section focuses on which properties need to be encoded in  $\varphi$  (genConstraints) to satisfy Rust semantics. For simplicity, in the remainder of this section, we assume that all API functions are not polymorphic. Section 5 describes in detail how we perform polymorphic API refinement (refine) and carry out hybrid instantiation to make APIs concrete.

## 4.2 Modeling Programs as SAT Formulas

The syntax of straight-line programs that are synthesized by SyRust is shown below. To simplify the synthesis process, we restrict the programs generated to not include branching and loops. However, our evaluation (Section 6) shows that even when only generating straight-line programs it is still possible to find complex memory bugs.

$$\begin{aligned} \text{Program} &:= \text{Line} \mid \text{Line}; \text{Program} \\ \text{Line} &:= f(\text{Vars}) \mid \text{let } v : \tau' = f(\text{Vars}) \\ \text{Vars} &:= v_1, \dots, v_k \end{aligned}$$

To build a SAT formula that represents the space of all possible programs of size  $l$ , we start by defining a set of Boolean variables that will be useful to model Rust semantics. These variables can be split into two main categories: *API variables* and *synthesis type-context variables*. API variables denote when an API function occurs in the program and synthesis type-context variables denote the universe of code variables that can be used in each line of the program.

For each API function  $f \in \mathcal{A}$  we create  $A_{f,i}$  Boolean variables with  $1 \leq i \leq l$ . We say that  $f$  is called on line  $i$  if  $A_{f,i}$  is set to *true* in  $\sigma$ . Additional restrictions are added to ensure that for each line  $i$  exactly one  $A_{f,i}$  is set to *true*.

To have valid programs, we must ensure that an API function  $f : \tau_1 \times \dots \times \tau_k \rightarrow \tau'$  can only be called on line  $i$  if the program contains variables that have been declared in previous lines with types  $\tau_1, \dots, \tau_k$ , respectively. We use a mapping from variable names ( $v$ ) to types ( $\tau$ ) in *synthesis type context*. For any line, there is a synthesis type context before it, which is the set of variables and types available before the function call, and a synthesis type context after it, which is the set of variables and types available after the function on this line has been called.

We model synthesis type contexts as subsets of the Cartesian product between the set of variable names ( $S$ ) and the set of types ( $T$ ).  $S$  is the union of variable names in the template and the new variable names that may appear in each line of code  $\{v_i \mid 1 \leq i \leq l\}$ .  $T$  is the union of types in the template, with the output types from the APIs  $\{\tau \mid f \in \mathcal{A}, f : * \rightarrow \tau\}$ . We define Boolean variables  $V_{v,\tau,i}$  with  $v \in S$ ,  $\tau \in T$ , and  $1 \leq i \leq m$ . We say that if variable  $v$  with type  $\tau$  occurs in the synthesis type context of line  $i$  then  $V_{v,\tau,i}$  is set to *true* in  $\sigma$ . Using these variables we can recursively construct the synthesis type context for each line. First, let us consider  $C_1$ , the synthesis type context for line 1. This is the synthesis type context before any APIs are called. This means the only variables available are those provided by  $\mathcal{T}$ .  $C_{i+1}$  must propagate every possible variable in  $C_i$  and include any new variables defined on line  $i$ . The following recurrence formalizes how we construct a synthesis type context for every line in the length- $l$  program.

**Definition 1** (Construction of Synthesis Type Context).

$$\begin{aligned} C_1 &= \{V_{x,\tau,1} \mid x : \tau \in \mathcal{T}\} \\ C_{i+1} &= \{V_{x,\tau,i+1} \mid V_{x,\tau,i} \in C_i\} \cup \\ &\quad \{V_{x_{i+1},\tau',i+1} \mid f \in \mathcal{A}, f : * \rightarrow \tau'\} \end{aligned}$$

We write  $C_l$  to denote the set of Boolean variables that occur in the synthesis type context on line  $l$ . We overload  $C_l$  as a mapping from variables to types, e.g., if  $V_{v,\tau,l} \in C_l$  and  $V_{v,\tau,l}=1$  then  $C_l(v) = \tau$ , and define the domain of the existing types of a context as  $\text{dom } C_l = \{d \mid \exists \tau \text{ st. } C_l(d) = \tau\}$ .



### 4.3 Basic Typing Constraints

We encode a set of rules that guarantees the correct usage of API functions over a set of variables. These rules are compatible with most imperative languages.

**Rule 1.** *If  $\mathcal{T}$  provides the inputs  $\{i_1:\tau_1, \dots, i_n:\tau_n\}$ , then  $\forall_{1 \leq j \leq n} C_1(i_j) = \tau_j$ .*

Rule 1 states that the available variables and types at the start of the synthesis are the ones in the template  $\mathcal{T}$ .

**Rule 2.** *Let  $Line_i = \text{let } x : \tau = f(\dots). C_{l+1}(x) = \tau$  and  $\forall v \in (dom C_l \cap dom C_{l+1}) \setminus \{x\}, C_{l+1}(v) = C_l(v)$ .*

Rule 2 describes how the synthesis type context is modified from line  $i$  to line  $i + 1$  with the call to function  $f$  in line  $i$ .

Next, we define a line  $l$  type checks with Context  $C_l$ , written  $C_l \vdash Line_l \Rightarrow C_{l+1}$  as follows.

**Definition 2** (Single-Line Type Check). *Let  $f : \tau_1 \times \dots \times \tau_k \rightarrow \tau'$ . A line  $l$  of code containing  $f(v_1, \dots, v_k)$  type checks with Context  $C_l$  if all of the following conditions are satisfied.*

1.  $v_1, \dots, v_k \in dom C_l$
2.  $\forall_{1 \leq p \leq k} C_l(v_p) \sqsubseteq \tau_p$
3.  $\forall_{1 \leq p < q \leq k} compatibleTypes(C_l, v_p, \tau_p, v_q, \tau_q)$

We use the subtype operator ( $\sqsubseteq$ ) instead of  $=$  to not miss generating valid programs, in particular for the following two aspects. First, Rust has subtyping for reference mutability:  $\&mut \text{String} \sqsubseteq \& \text{String}$ . Second, by allowing type variables to match the broadest range of types (i.e.,  $\forall \tau, \tau' \sqsubseteq T$ ), the synthesis can generate code that uses polymorphic types; for instance  $Vec\langle String \rangle \sqsubseteq Vec\langle T \rangle$ . However, if the same type variable occurs in multiple places, then it must be matched to compatible types. For example, for  $Vec :: \text{push}$ , we must ensure  $T$  from first input  $\&mut Vec\langle T \rangle$  matches with a type that is compatible with what  $T$  matches in the second input. To this end, we use a  $compatibleTypes$  function that determines if a polymorphic match is compatible for every pair of polymorphic variables  $v_i, v_j$ . Note that if  $v_i$  and  $v_j$  are not polymorphic then they are always compatible. More discussions on polymorphic types are presented in Section 5.

We can inductively extend Definition 2 to define type checking for a multi-line program below.

**Rule 3.** *Let  $Program = Line_0; \dots; Line_{n-1}$  and  $Line_i$  contain a call to  $f_i$ . Then there exists  $C_1, \dots, C_n$  such that for every  $1 \leq i \leq n, C_i \vdash Line_i \Rightarrow C_{i+1}$ .*

### 4.4 Ownership and Variable Lifetime

We show how to encode variable and reference lifetime constraints in Rust as described in Section 2.

**4.4.1 Basic variable lifetimes.** A variable  $v$ 's lifetime, denoted  $L_v$ , is the set of *continuous* contexts in which  $v$  occurs:

**Definition 3** (Variable Lifetime). *A lifetime for variable  $v$  that is created on line  $s$  and ends in line  $e$  is defined as follows:  $L_v = \{C_l \mid s \leq l < e, v \in dom C_l\}$ .*

Some types can be used at most once in a line. The algorithm needs to add the following condition to the  $compatibleType$  function presented in Definition 2. If the variables are the same, then the types must be a primitive type or static reference ( $\&$ ), the only type of reference that is allowed to be used in interleaving terms on the same line. That is,

**Rule 4.**  *$compatibleTypes(C, v, \tau_i, v, \tau_j)$  with  $i \neq j$  is true if and only if  $\tau_i = \tau_j$  is a static reference.*

Further, non-primitive variable lifetimes need to be terminated upon use.

**Rule 5.** *Let  $v$  be a non-primitive variable used on the right-hand-side of  $Line_i$ , then  $v$ 's lifetime terminates here. We require that  $v \notin dom C_{i+1}$ , and therefore  $C_{i+1} \notin L_v$ .*

**4.4.2 Relationship Between Lifetimes.** Next, we encode rules for references. In particular, we need to define the notion of a variable outliving another one.

**Definition 4** (Lifetime Containment). *A lifetime  $L_{v_1}$  contains another lifetime  $L_{v_2}$  if and only if  $L_{v_2} \subseteq L_{v_1}$ .*

Because our algorithm treats borrowing ( $\&$  and  $\&mut$ ) as a special kind of API, the lifetime containment rule will be formalized as follows.

**Rule 6.** *Let  $Line_i$  be a call ( $\&$  or  $\&mut$ ) to borrow variable  $v$  to produce a reference  $v'$ , which points to  $v$ , then  $L_{v'} \subseteq L_v$ .*

**4.4.3 Paths and Lifetimes.** Rule 6 alone does not capture all of the constraints in lifetime management; constraints on implicit movements of references are missing. The Rust type signatures are annotated with lifetime variables [4, 6] (e.g. 'a). These annotations can be programmer-provided, or in some cases, the Rust compiler automatically annotates using Lifetime Elision rules [4].

Consider the option type  $Option\langle T \rangle$ . Assuming it is not  $None$ , it serves as a wrapper for that can be unwrapped with  $unwrap : Option\langle T \rangle \rightarrow T$ . This carries no significance when  $T$  is not a reference. However, when  $T$  is a reference, we need to make sure that the output of type  $T = \&mut \tau$  must be as a valid reference and all of the reference rules must be applied to it. To formalize implicit movements, we define *Path* to model data flow.

**Definition 5.** *Path A path is a sequence of variable-line pairs  $(v_s, l_s), \dots, (v_k, l_e)$  such that for every  $(v_i, l_j)$  the variable associated with  $l_{j+1}$  is not equal to  $v_i$ . Moreover, there exists a function call on  $Line_j$  that propagates the lifetime of  $v_j$  to the new variable defined on  $Line_j$ .*

This rule formalizes the notion that we need to keep the lifetime relation even when the types change and content is moved out.

**Rule 7.** *For every path  $(v_s, l_s), \dots, (v_k, l_e)$ ,  $L_s \subseteq L_{v_i}$  is enforced for every  $v_i$  in the path.*

Note that the above rule is enforced using a post-processing check, rather than encoded in the formula. This is because this notion of connectivity may lead to a cubic number of additional SAT clauses [26] and it is more efficient to either solve it lazily or via a post-processing check.

**4.4.4 Banned Operations and Lifetimes.** Recall the example in Section 2 where creating a mutable reference (vr2) was prevented because another mutable reference (vr) was already active in the context. Our goal is to add constraints to prevent these kinds of programs from being synthesized. These constraints fall into the following three categories:

1. while  $\&\text{mut}$  is active, prevent another  $\&\text{mut}$ ,
2. while  $\&\text{mut}$  is active, prevent another  $\&$ ,
3. while *any*  $\&$  is active, prevent another  $\&\text{mut}$ .

In the last category, *any* is emphasized as it is possible to have multiple static references, but every static reference must be exclusive against mutable references.

**Rule 8.** *If  $\text{Line}_i$  is a call to mutable borrow ( $\&\text{mut}$ )  $v$  and the resulting mutable reference is stored in  $v'$ , then  $\forall \text{Line}_j$  such that  $j > i$  and  $v' \in \text{dom } C_{j+1}$ ,  $\text{Line}_j$  cannot be a mutable borrow or static borrow of  $v$ .*

**Rule 9.** *If  $\text{Line}_i$  is a call to static borrow  $\& v$  and the resulting mutable reference is stored in  $v'$ , then  $\forall \text{Line}_j$  such that  $j > i$  and  $v' \in \text{dom } C_{j+1}$ ,  $\text{Line}_j$  cannot be a mutable borrow of  $v$ .*

## 4.5 Connecting to SAT

We have presented the rules that we use to encode the semantics of Rust in our synthesis engine. We convert these rules into a Boolean Satisfiability formula and solve it using an off-the-shelf SAT solver. We refer the reader to Appendix A for further details on the SAT encodings.

Getting back to Algorithm 1, the rules of this section correspond to what genConstraints generates. Once this formula is solved, we will have a truth assignment  $\sigma$  for the Boolean variables that correspond to a valid program in Rust. Using the one-to-one correspondence between the Boolean variables that represent the APIs and inputs with their code representation, the codeGen function walks the solution and builds the test case line-by-line.

## 4.6 Properties of the Synthesized Programs

**Remark 1.** *The programs synthesized by SyRust satisfy the requirements of the Rust compiler for the following properties.*

1. *The synthesized program type-checks (Rules 1, 2, 3).*
2. *Variable aliases are not allowed except for static references (Rule 4).*
3. *All non-primitive variables are uniquely owned (Rule 5).*
4. *While any references are active, the memory location cannot be moved or deallocated (Rules 6, 7).*

5. *For any given variable, there is at most 1 active mutable reference or an arbitrary number of active static references, but not both (Rules 8, 9).*

Even though the synthesized programs can still result in compilation errors, our evaluation (Section 6) shows that only a very small percentage of programs are rejected by the compiler and these errors are due to other factors not captured in the above properties, e.g., polymorphic errors that can be lazily fixed with the approach presented in Section 5.

## 5 Hybrid Refinement

In this section, we present our *hybrid* approach to instantiate polymorphic APIs. Rust allows users to specify that a type variable can only be instantiated with types implementing a given trait (i.e., types support a set of functions). We also discuss how SyRust deals with traits here.

A type  $\tau$  is concrete if it has no type variables, and polymorphic otherwise. APIs with only concrete inputs and output are concrete, and otherwise polymorphic. The following are the three most common types of polymorphic Rust APIs

1. No Input Polymorphism:  
*Example:*  $\text{Vec} :: \text{new} () \rightarrow \text{Vec}\langle T \rangle$
2. Polymorphic Input, Concrete Output:  
*Example:*  $\text{Vec} :: \text{push} (\&\text{mut Vec}\langle T \rangle, T) \rightarrow ()$
3. Polymorphic Input, Polymorphic Output:  
*Example:*  $\text{Vec} :: \text{pop} (\&\text{mut Vec}\langle T \rangle) \rightarrow \text{Option}\langle T \rangle$

While these categories are not exhaustive; it is certainly possible to write a function that has type variables in the inputs different from type variables in the output. In practice, we find this extremely uncommon.

Our observation is that neither purely lazy nor purely eager approaches to polymorphism are sufficient for Rust. Purely lazy approaches cannot synthesize types for no input polymorphism, and purely eager approaches result in too many incorrect APIs (see Section 7.3 for more details). Therefore, SyRust combines the two.

### 5.1 No Input Polymorphism

Let's start with the most simple case: a function with a polymorphic output and no inputs. These functions are often used as constructors in data structure libraries. Since the function has no inputs, we cannot infer the concrete output type automatically. The Rust compiler is often able to reason about the output by looking ahead (i.e. Resolve to  $\text{Vec}\langle i32 \rangle$  because the vector gets  $i32$  pushed onto it sometime later). Such information is not available to our synthesis algorithm that builds programs from the ground up. Instead, we opt to eagerly concretize the output type. We need to find concrete types to substitute type variables with. As we noted in Section 2, we must be careful as this space is infinitely large. The hybrid refinement module mines concrete types from the API set and template. For example, if we collect  $i32$  and  $u32$  from the API and template, we add 2 variants of  $\text{Vec} :: \text{new}$

where the output type is changed to `Vec<i32>` and `Vec<u32>` respectively. This combinatorial enumeration generates a large number of APIs, so it should be only sparingly used. We only use it for No Input Polymorphism.

The eager concretization also ignores trait annotations on type variables. This means some of the type instantiations are incorrect due to trait mismatches. Because the instantiated APIs are concrete, they cannot be refined further. SyRust removes fully concrete APIs that caused trait errors to avoid a combinatorial explosion of incorrect programs.

## 5.2 Polymorphic Inputs, Concrete Output

Let us now consider cases where a function has polymorphic inputs and concrete output. We only need to instantiate the inputs. For a majority of the cases, the subtyping strategy outlined in Section 4 work fine. However, this strategy fails when type variables in the input are annotated with traits that any matching type must support.

While it is certainly possible to support traits at the synthesis level, Rust's trait system is fairly complex. Traits can be polymorphic themselves, and a few complex traits are even defined on other traits. Instead of dealing with complex trait requirements, we use the compiler errors as feedback to refine polymorphic variables. When a type match fails because of mismatching traits, we refine the type variable with metadata to no longer match that particular type. The refinement is usually complete after a few rounds.

## 5.3 Polymorphic Inputs, Polymorphic Output

In the final category, let's consider, `Vec :: push : &mut Vec<T> → Option<T>`, which takes a mutable reference of a vector and returns a wrapped value containing the vector element last and nil if the vector is empty. Because this output may be used in later API calls, we must get the exact type of the output. We must do so without modifying the original API as we want to use to match other types later. Our solution is to duplicate the function for that set of inputs only.

We begin the process the same way as we deal with functions that have no polymorphic outputs. For example, the synthesis algorithm matches `&mut Vec<T>` with a concrete type such as `&mut Vec<i32>`. The synthesizer generates a test case where `Vec :: push` is used with the input of type `&mut Vec<i32>` and predicts the output type as `Option<i32>`. This test case is then passed to the test executioner, which attempts to compile and run the test case.

If the test case compiles successfully, we duplicate the function and fully concretize the duplicated API's inputs with the current input types. Since the program type-checked, the predicted output type is correct. We replace the output of the duplicated API with the prediction.

If the test case fails to compile, what happens next depends on what error we get. If the error is originating from the inputs to the API, then we perform the same refinement as in the previous section. Other errors are fixed directly.

For example, if we get the message “expected String, got `Vec<i32>`” then we duplicate the API, concretize the inputs, and set the output type to `Vec<i32>`.

Because we are duplicating APIs, we must be careful to keep the duplicated API disjoint from the original. After duplicating the API, we add metadata to the original to prevent it from being used with the same combination of input types. We block combinations rather than individual inputs because APIs like Hash Map's get function has multiple type variables involved, and blocking individual inputs will block too many potential candidates.

## 6 Implementation and Evaluation

We implemented our proposed approaches in SyRust, an automatic testing tool leveraging semantic-aware synthesis and type refinement. We evaluated it against Rust libraries on the following research questions.

- RQ1.** How effective is SyRust in generating valid Rust test cases and in finding bugs in Rust libraries?
- RQ2.** How do the various semantic-aware synthesis features (Section 4.4) help generate valid test cases?
- RQ3.** How does the polymorphic type refinement algorithm help find bugs?

### 6.1 Implementation

The synthesizer is implemented as a Java program that takes in a configuration file that specifies the APIs to be tested and a template to use in the tests, and generates test cases through constraint solving. We use Sat4J [10] to solve SAT formulae. The type refinement is also implemented in Java.

In addition to using the constraints in Section 4, the synthesizer encode rules to prune redundant programs. For example, a program is redundant if it contains a borrow operation but the reference created by that borrow is never used. Since borrowing does not alter the state of the data, we know this program is equivalent to a shorter program with this borrow operation removed. The synthesizer block such patterns through the rules formalized in Appendix B.

The test executioner uses cargo, the official build system for Rust. We run cargo with the `-message-format=json` flag to get compiler errors in JSON format, and send back the parsed data to the synthesizer. To detect bugs, we use Miri [22], a bug detection interpreter already integrated with cargo. Miri relies on the Rust compiler to generate the Mid-Level Intermediate Representation (MIR), which Miri then interprets. It flags any behavior that is considered unbounded according to its dynamic semantic model for Rust [19].

The API type signatures used for synthesis are collected using a modified Rust compiler. We compile the target libraries with the modified compiler and derive the API type signatures. Since the total number of APIs is too large, a smaller subset is used in testing. Section 6.2 provides further details on how this selection is done for our experiments.



## 6.2 Library and API Selection

To select a sufficiently large and representative set of Rust libraries for our evaluation, we examine `crates.io`, the official library repository for Rust. We focus on Data Structure and Encoding categories. This is because these libraries are often used as the building blocks of other libraries and thus are highly critical to the Rust ecosystem. Furthermore, they are also more likely to contain unsafe code as they often perform low-level operations. From the two categories, we select the 30 libraries by download count, while making sure that the libraries have the following 2 properties.

1. **The library and all of its dependencies are written purely in Rust.** This is required as Miri cannot handle foreign functions.
2. **The library must also be API based.** Libraries consisting mostly of macros cannot be used as no type signatures are readily available for these libraries.

We also prioritize libraries that contain unsafe code. For details about library versions, popularity, and other features, we refer the reader to Appendix C, Figure 11.

For evaluation, we allow a small number of specific APIs to be manually selected to simulate the scenarios where the programmers want to test specific APIs. In our experiments, we selected no more than 2 APIs per group by this method. Then the rest of the APIs are selected using a combination of heuristics (e.g., prioritizing constructors) and random selection. A total of 15 APIs, including the 2 specially selected ones, per library is used in the testing. The set of 15 is appended with 3 default APIs for mut casting (`assign`) and borrowing (`borrow`, `borrow_mut`). In practice, we find that this number of APIs gives us a good trade-off between having a diverse set of APIs and synthesizing complex programs with multiple lines of code.

## 6.3 Experiment Setup

Since SyRust generates a large number of test cases, the majority of the computational resources is spent on running these test cases. Exploiting the fact that the test cases can be run independent of each other, we deployed 64 containers running the test executioner across 4 machines, each of them running an Intel i9-9900K with 128GB of RAM. On the other hand, we find that in practice, solving the constraint formulas is quite fast. Therefore, the synthesizer node is deployed on an Intel i7-6700K with 32GB of RAM. All nodes are running Ubuntu Linux, with the test execution nodes running 19.10 and the synthesizer running 18.04. We used the 2020-10-01 nightly build of Rust and Miri (Miri only comes with nightly builds). Nodes communicate using ZeroMQ, with the synthesizer node running using JeroMQ Java library version 0.5.0 and test executioners using PyZMQ version 19.0.0 as the respective driver. Due to memory constraints, we limit the size of the queue at 10,000 programs.

## 7 Evaluation Results

We present our evaluation results and answer RQ1-RQ3 here. Due to the lack of prior work in Rust library API testing, we are unable to provide an external baseline.

### 7.1 RQ1: New Bugs and Test Effectiveness

We ran SyRust on 30 libraries with a timeout of 10 hours. We removed `cookie-factory` and `jsonrpc-client-core` as they use first-order functions, not supported by our syntax. For libraries that overlap (e.g. `crossbeam-queue` and `crossbeam`), we test different components with no shared APIs. The results are summarized in Figure 4. The columns denote the library name, number of lines enumerated, number of test cases processed, the number of programs that were rejected by the compiler, followed by a percent-wise breakdown of the rejections into categories. The “Type errors” are caused by wrongly instantiated polymorphic types; “Lifetime & Ownership” is self-explanatory; “Miscellaneous errors” typically indicate errors in collected API type signatures (done in the experiment setup phase). Finally, Libraries that were flagged as buggy by SyRust are marked with a ★.

Overall, we observe that for most libraries, the portion of test cases rejected by the compiler is extremely low (often less than one percent) indicating the effectiveness of our tool in generating valid Rust programs. However, for certain libraries, the error rate is significantly higher than for the rest and this is due to features that are not supported yet in SyRust. These include an unsupported lifetime corner case involving anonymous parameterized lifetimes in `slab`, `sval`, `csv-core`, and a few others libraries, missing default values for type variables in `Petgraph`. We believe the former issue can be fixed with improved API collection and fixing the latter requires modifying the rules of Section 5 and Section 4 to accommodate default values of type variables. We leave the implementation to future work.

Miscellaneous errors, such as “expected n arguments, found j,” are often caused by errors in API type signatures. However, the large number of miscellaneous errors in `generic-array` and `hashbrown` are “method not found” errors. Notably, this error can be caused by both polymorphism mistakes and API issues. We chose to conservatively include them in Miscellaneous, but we suspect a large number of them are polymorphism induced.

Finally, we observe that some libraries have significantly fewer test cases synthesized than others. For most such libraries, there are simply not enough valid combinations to generate a large number of valid test cases. This means that the synthesis terminates early, with a small number of solutions found. These libraries also exhibit higher error rates as APIs become more refined towards the end of the run. One exception is `dashmap`, where the library was extremely slow to be interpreted by Miri. Therefore, only about half as many test cases can be executed within the 10 hour limit.



Library	Max Test Case Length	# Synthesized	# Rejected (% of total)	Type Error (%)	Lifetime & Ownership (%)	Misc. (%)
smallvec	9	1225952	66 (< 0.01 %)	95.45 %	0.00 %	4.55 %
crossbeam-utils	5	1242990	703 (0.06 %)	58.61 %	2.99 %	38.41 %
bytes	7	1194703	43 (< 0.01 %)	93.02 %	0.00 %	6.98 %
slab	6	1229924	811 (0.07 %)	64.24 %	35.76 %	0.00 %
crossbeam-deque	6	1216076	957 (0.08 %)	100.00 %	0.00 %	0.00 %
generic-array	10	1216241	2088 (0.17 %)	1.29 %	0.00 %	98.71 %
crossbeam-queue ★1	5	1153651	19463 (1.69 %)	100.00 %	0.00 %	0.00 %
num-rational	4	1255485	4506 (0.36 %)	99.82 %	0.00 %	0.18 %
hashbrown	6	1122196	17649 (1.57 %)	7.25 %	0.54 %	92.21 %
crossbeam ★2	4	1231844	1367 (0.11 %)	98.76 %	0.88 %	0.37 %
petgraph	4	1318138	143347 (10.87 %)	100.00 %	0.00 %	0.00 %
im-rc	6	1226829	25531 (2.08 %)	95.06 %	0.06 %	4.87 %
bitvec ★3	7	1221730	120 (< 0.01 %)	100.00 %	0.00 %	0.00 %
ndarray	9	1188730	830 (0.07 %)	100.00 %	0.00 %	0.00 %
dashmap	7	660986	918 (0.14 %)	82.57 %	0.22 %	17.21 %
encoding_rs ★4	6	1233420	152 (0.01 %)	100.00 %	0.00 %	0.00 %
bstr	9	1207815	258 (0.02 %)	93.80 %	1.55 %	4.65 %
csv-core	6	14961	478 (3.19 %)	5.02 %	93.72 %	1.26 %
data-encoding	10	900509	136 (0.02 %)	64.71 %	23.53 %	11.76 %
encode_unicode	6	1238800	67 (< 0.01 %)	94.03 %	0.00 %	5.97 %
urlencoding	6	1139257	48 (< 0.01 %)	100.00 %	0.00 %	0.00 %
rmp-serde	6	11544	963 (8.34 %)	99.27 %	0.00 %	0.73 %
bytemuck	5	112030	19568 (17.47 %)	86.26 %	13.74 %	0.00 %
sval	10	86606	392 (0.45 %)	44.39 %	55.61 %	0.00 %
base16	6	1194409	78 (< 0.01 %)	100.00 %	0.00 %	0.00 %
cbor-codec	6	17292	656 (3.79 %)	36.59 %	63.41 %	0.00 %
hcid	5	1158079	100 (< 0.01 %)	100.00 %	0.00 %	0.00 %
utf8-width	4	1267697	168 (0.01 %)	100.00 %	0.00 %	0.00 %

**Figure 4.** This table gives rejection rates for all the tested libraries, and breaks down the rejections by categories (right 3 columns). For every row, the right 3 cells add up to 100 %, showing which kind of error was dominant in the compiler rejection. Libraries in which we found bugs are marked with a ★. For information about each bug, see Figure 5.

In total, SyRust found 4 previously unknown bugs in 3 different libraries, all of which were accepted by the library authors. These bugs are shown in Figure 5. The columns denote the library in which the bugs were found, the minimum number of lines required to trigger the bug, the type of bugs, and whether the authors accepted the bug. We will go over these noteworthy bugs in detail.

★	Bug Type	Min. Lines to Induce	Time to Discovery (s)	Accepted by Authors?
★1	Memory Leak	1	4.45	Yes
★2	Hanging Pointer	3	2850.3	Yes
★3	Use-After-Free	5	200.61	Yes
★4	OOB Pointer	4	238.99	Yes

**Figure 5.** Bugs Caught by SyRust

Bug ★1 is a memory leak in crossbeam-queue’s ArrayQueue data structure that can be induced by initializing the

ArrayQueue with a non-zero initial capacity. While this is a one-line trivial bug, it exposes problematic assumptions about internal memory layout, and the issue was cited in a vulnerability report [3].

Bug ★3 is a dereference-after-free bug in bitvec, a bitvector library. This bug occurs when deallocating a non-empty BitBox object, a fixed-length equivalent of BitVec.

```

let x1_l1: BitVec<Msb0, usize> = BitVec::with_capacity(0);
let mut x2_0: = x1_l1;
let x3_0: &mut = &mut x2_0;
x3_0.push(true);
let x5_0: = x2_0.into_boxed_bitslice();

```

**Figure 6.** Buggy Case for bitvec

Looking at the bug-inducing code in Figure 6, we see that it is particularly challenging to synthesize. First, note that it involves ownership movement through the `into_boxed_bit_slice` function. This means that the same call sequence cannot be triggered using a loop-based fuzzing harness because it would not pass compiler checks when the ownership moves inside the fuzz loop. Therefore, it can only be triggered using synthesis-driven approaches like ours. Second, the `bitvec` library uses extensive polymorphism and trait-driven programming. For example, the type `BitVec<Msb0, usize>` is really `BitVec<0, T>` with `0` and `T` instantiated with `Msb0` and `usize`. One must be careful in the instantiation because `BitVec<usize, Msb0>` will not pass compiler checks, since `usize` does not support the `BitOrder` trait (`Msb0` = Most significant bit is index 0). Finally, the sheer size of the bug-inducing case makes this a difficult bug to induce.

Finally, bugs `★2` and `★4` provide interesting insight into the operational semantics of Rust. These bugs may look benign, as simply having a out-of-bounds (OOB) or hanging pointer is not considered to be harmful as long as it is not dereferenced. However, `Miri` considers it to be buggy, because Rust's standard library provides the `std::mem::MaybeUninit` wrappers for dealing with pointers that may possibly be uninitialized. By not using `MaybeUninit`, the library risks becoming buggy if the compiler behavior changes at a later time. For example, if deallocation scheduling changes, it may cause a use-after-free bug because the compiler is not aware that this buggy pointer is pointing to.

Further discussions and the bug-inducing code are presented in Appendix C.

## 7.2 RQ2: Semantic Awareness

For this question, we examine how the SyRust's awareness of ownership and lifetime checks help generate valid test cases. Using 2 of the 4 libraries that had previously unknown bugs, we attempt to replicate the results with the Rust-specific semantic awareness constraints (Section 4.4) turned off. We excluded `crossbeam-queue` (`★1`) and `encoding_rs` (`★4`) because the former is a 1-line bug, and the latter is extremely simple with respect to lifetime/ownership. Once again, we run every experiment for 10 hours.

Before we dive into the setup, we note that these constraints are highly dependent on each other. Turning off one will cause the rest of the constraints to behave incorrectly (Appendix A). We believe these constraints should be evaluated as a whole, leading us to employ this on/off comparisons instead of fine-grained feature-for-feature comparisons.

Results are shown in Figure 7. The first row of graphs plot the growth of errors over the execution. The x-axis is time, from 0 to 10 hours. The y-axis is the percent of test cases up to that point that were rejected by the compiler, shown in logarithmic scale. This metric is cumulative, and the overall rejection rate is the y value of the rightmost point on the curve. Lower is better as it implies a low rejection rate. The

dotted red line is SyRust with semantic awareness turned off, and the solid blue is the baseline with all features.

Both `crossbeam` and `bitvec` see an increase in errors. However, the increase in error rate is far more severe for `bitvec`. This occurs because `bitvec`'s APIs have more APIs that move ownership or extend lifetime.

The second group is a breakdown of the errors by category, and across time. These graphs are only for executions with the semantic awareness turned off. The x-axis is again time, and the y-axis is the percent this error category occupies out of the total programs rejected by the compiler (in linear scale). The blue solid line corresponds to type errors, green dashed line corresponds to Lifetime/Ownership errors, and the orange dotted line corresponds to miscellaneous. This means the sum of the 3 lines is always 100%. Higher the line, more dominant this particular error is.

For both libraries, rejections rate due to lifetime and ownership errors rise to overtake type errors as the most dominant kind of error. Like with the previous case, this reversal in dominant errors type is quicker with `bitvec`.

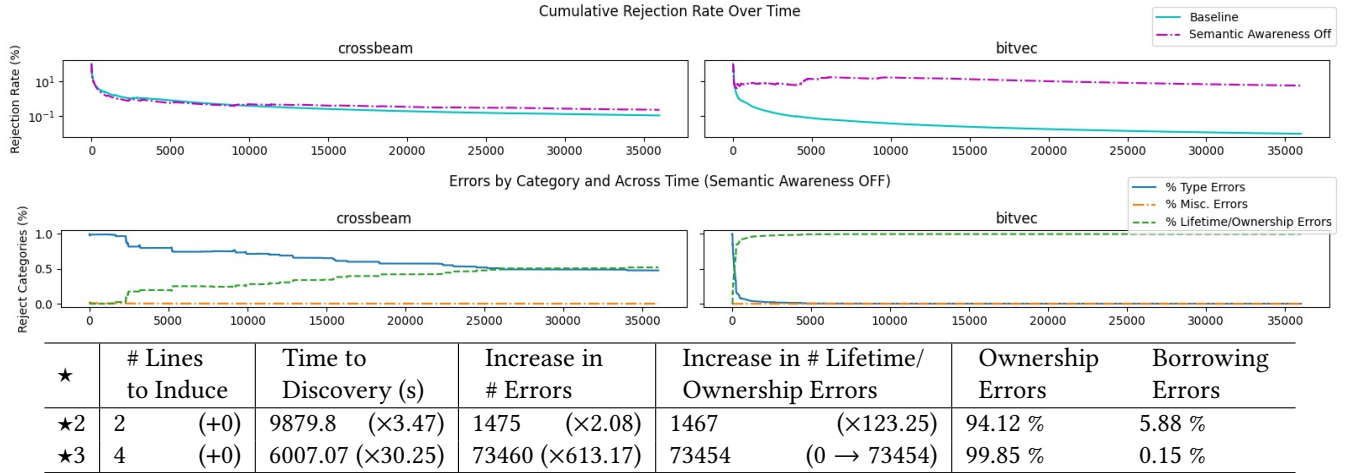
We also provide a table showing the time to bug discovery, number of general errors, number errors for lifetime/ownership, and the breakdown of lifetime/ownership errors by sub-category. In parenthesis, we show the increase over the baseline. As expected, we see a significant increase in compiler rejections and time to discovery for both libraries. Again, the negative effects are more prominent for `bitvec`. Meanwhile, the length of the bug-inducing case is the same. Finally, the vast majority of the Rust-specific rejections are ownership errors (Rules 4-7 of Section 4), with only few percent are lifetime errors (Rules 8 and 9).

Our semantic aware synthesis significantly reduces the error rate and decreases the time to discover bugs. Increase in error rate is more prominent for complex APIs, making semantic awareness critical to scalability.

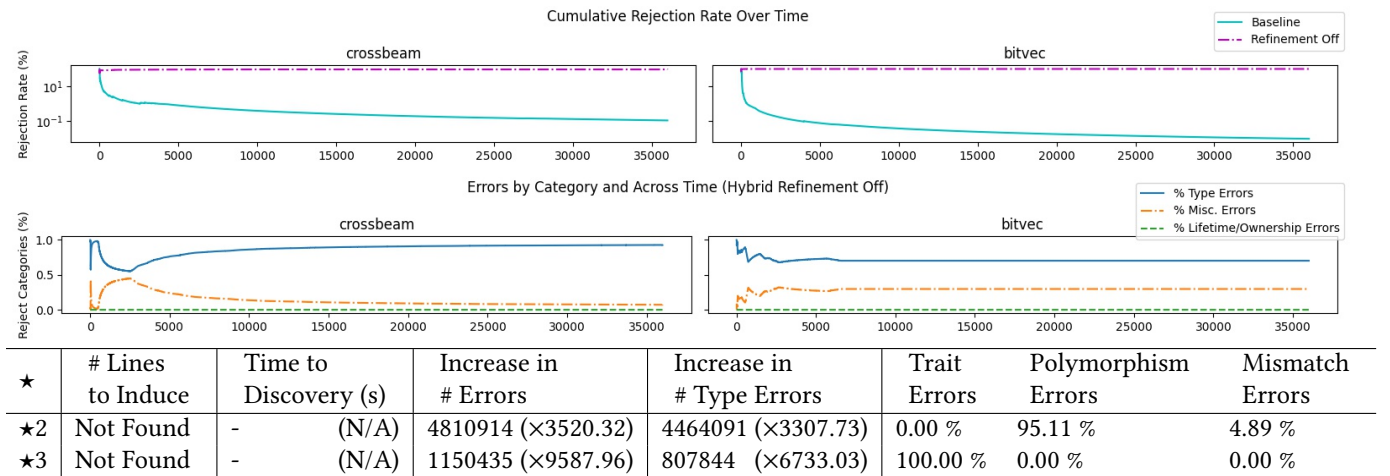
## 7.3 RQ3: Polymorphic Type Refinement

We measure how the incremental refinement procedure contributes to SyRust. However, unlike RQ2, simply turning this feature off will not work. Since the APIs are polymorphic, the type variables have to be filled with something to type check. One possible option is to naively apply the eager instantiation strategy used for no-input polymorphism (Section 5.1) to every API regardless of the category and turn off refinement. This eager strategy, which has also been used in past synthesis projects [13], will serve as the baseline.

The results are given in Figure 8. Since the targets and statistics we collect are identical to those for RQ2, the axes of the the graphs we generate are also identical. On the other hand, we are now concerned with polymorphism and provide breakdowns of type errors instead of lifetime errors. Since the bugs were unreachable with polymorphism turned off, the left two columns are empty.



**Figure 7.** Selected libraries with semantic awareness turned off. The top row shows error rates regardless of kind. Blue is baseline and purple is with semantic awareness turned off. Second row of graphs show the portion occupied by each error kind. Green is Lifetime/Ownership, blue is type error, and orange is misc.



**Figure 8.** Selected libraries with hybrid API refinement turned off. The top row shows error rates regardless of kind. Blue is baseline and purple is with refinement turned off. Second row of graphs show the portion occupied by each error kind. Green is Lifetime/Ownership, blue is type error, and orange is Misc.

The top 2 graphs show a clear increase in the rejection rate, while no-refinement for `bitvec` and `crossbeam` leads to rejection rates of near 100%. For both libraries, type errors are the most dominant. Type errors start out being the most dominant, falter in the first few hours, but stabilize and maintain dominance after some time. We also see a rise in miscellaneous errors. Because these errors occur mostly due to API collection errors (see RQ1), we believe this error was compounded when the eager type instantiation algorithm instantiated type variables with concrete types that had incorrect import paths. Finally, we note an unsurprisingly large increase in both overall rejection rates and those limited to type errors. The most dominant kind of error was different

between the two libraries. `Bitvec`'s type errors were entirely due to traits, and `crossbeam` suffered from polymorphism errors with some cases of outright type mismatch.

We showed that the hybrid polymorphic refinement is critical to the success of SyRust. In particular, it allows SyRust to scale to complex polymorphic APIs, and find bugs that cannot be reached without our hybrid approach.

#### 7.4 Limitations and Future Work

While our tool has been shown empirically to be capable of finding bugs, some limitations remain. A few of the limitations are inherited from the Rust tooling, and we expect them to be fixed as the Rust ecosystem evolves. Others are



more fundamental limitations, and offer insight into a need for further work in this area. We discuss them below.

**7.4.1 Closures and Asynchronous Functions.** Recall the limitation from RQ1 where SyRust failed to generate valid test cases for `cookie-factory` and `jsonrpc-client-core`. This is due to lack of syntax support for closures. Closures in Rust are used to define anonymous functions that are treated as first-class objects. Because our synthesizer is limited to straight line code, defining new closures is not allowed. Because asynchronous functions rely on closures to perform callbacks, asynchronous APIs are off the table as well. Given some recent work in synthesis of asynchronous programs [9], we believe this will be an interesting area of research.

**7.4.2 Inputs to the Test Program.** Currently, our test framework runs on user-provided input parameters, and we do not mutate them. Mutation of the inputs is likely to trigger more bugs, but we chose to focus on the orthogonal API problem for now. Because this problem may be addressed with Bounded Model Checking (BMC) techniques, we believe that extending our tool into a BMC tool in the style of CRUST [27] is an interesting path forward.

**7.4.3 Optimal Scheduling of Tests.** We run the tests in whichever order the SAT formula was solved. This is not necessarily the best order, and we may be able to find bugs quicker if we prioritize certain programs. This is a well studied problem in the fuzzing community [30] and we expect that some of those techniques are applicable to our domain.

**7.4.4 Miri Limitations.** As noted in Section 6, we use Miri [22] to run our test cases. While Miri will not miss any bugs, its precision comes at the cost of speed. Given all the checks performed, Miri is significantly slower than running the compiled code. In particular, some of the Stacked Borrows rules [19] results in cases where the loops of linear complexity require polynomial number of operations to interpret. This significantly limits our ability to test APIs that internally contain large loops and iterations.

## 8 Related Work

Automatic test case generation is useful for testing language interpreters and compilers [12, 15, 16, 32] and libraries [7, 20, 21, 23]. Some mutate and pieces together existing code segment [15, 18], some generate code from scratch [12, 32], and some like ours, generate API call chains [7, 23].

A number of techniques have been proposed for API testing. Tools like RESTler [7] mutate call sequences [33]. FUDGE [8] and FuzzGen [17] generate fuzz drivers for C/C++ libraries from existing code bases. FUDGE uses recursive search to combine API calls. FuzzGen performs program analysis of API usage to generate A<sup>2</sup>DG graphs that represent data and control dependencies and uses the graph

to generates programs. Since these works are for C/C++, lifetime and ownership is not an issue.

Leveraging advances in constraint-solving, component-based program synthesis techniques [18] have been applied to many languages. Most closely related to ours, are SyPet [13] (Java) and H+ [14] (Haskell), which use constraint-solving-based technique to generate loop-free straight-line code. Both SyPet and H+ employ a graph-based encoding, which is unwieldy for encoding the lifetime and ownership of variables; and therefore SyRust does not use. To handle polymorphism, SyPet takes a completely eager approach to Java Generic Types. To support Haskell's polymorphism and partial evaluation, H+ starts by assuming every API takes and returns the universal polymorphic type and refines lazily from there. SyRust takes a hybrid approach. Since Rust has polymorphic struct constructors but not partial evaluation, we start with a mix of both concrete and polymorphic types and use the hybrid approach to refinement.

Dewey et al. [12] encode the Rust syntax and semantic constraint in Prolog [11] to generate code to test the Borrow Checker. Their lifetime and ownership constraints are roughly equivalent to ours. They generate programs with complex structures such as loop nesting and define and use of data structures. Because they aim to test the compiler, they only support polymorphism through a few standard library features like `Box<T>` and cannot be easily modified for API testing. Their tool is more resource-intensive than ours, owing to the use of Prolog.

Several projects aim to formalize Rust's operational semantics [25, 28], Stacked Borrows [19], being the most recent. Stacked Borrows is used as a correctness model in Miri [22], a sub-component in SyRust to detect unsound behaviors.

Finally, Oxide [29] formalizes static checks of Rust. We could connect to it to formally verify the completeness of our constraints as future work.

## 9 Conclusion

We have proposed a semantic-aware synthesis algorithm for Rust that effectively synthesizes valid API usage across a wide range of libraries exercising complex language features like ownership and polymorphism. Our algorithm encapsulates Rust's compiler checks through a logical encoding and leverages a hybrid API refinement strategy to drive polymorphic APIs. We implement our contributions in SyRust, a library testing framework for Rust. Our experiments demonstrate SyRust's ability to generate valid test cases for a wide range of libraries and find 4 confirmed bugs. In future work we plan to investigate whether coverage information can be used to improve the synthesis of valid test cases. We also plan to apply our framework to testing the Rust compiler and to formally prove the correctness of the synthesis.

## References

- [1] 2018. *Built In Rust*. <https://www.rust-lang.org/what/>
- [2] 2018. *CVE-2018-1000810*. <https://nvd.nist.gov/vuln/detail/CVE-2018-1000810>
- [3] 2020. *Anonymized for double-blind review*. <https://github.com/RustSec/advisory-db/XXX>
- [4] 2020. *Built In Rust*. <https://doc.rust-lang.org/nomicon/lifetime-elision.html>
- [5] 2020. *Rust's Unsafe Code Guidelines Reference*. <https://rust-lang.github.io/unsafe-code-guidelines/>
- [6] 2020. *Trait and lifetime bounds*. <https://doc.rust-lang.org/reference/trait-bounds.html>
- [7] Vaggelis Atlidakis, Patrice Godefroid, and Marina Polishchuk. 2019. RESTler: Stateful REST API Fuzzing. In *Proceedings of the International Conference on Software Engineering*. IEEE Press, 748–758.
- [8] Domagoj Babic, Stefan Bucur, Yaohui Chen, Franjo Ivancic, Tim King, Markus Kusano, Caroline Lemieux, László Szekeres, and Wei Wang. 2019. FUDGE: Fuzz Driver Generation at Scale. In *Proceedings of the ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering*. ACM, 975–985.
- [9] Suguman Bansal, Kedar S. Namjoshi, and Yaniv Sa'ar. 2018. Synthesis of Asynchronous Reactive Programs from Temporal Specifications. In *Proceedings of the International Conference on Computer Aided Verification*. Springer International Publishing, 367–385.
- [10] D. L. Berre and Anne Parrain. 2010. The Sat4j library, release 2.2. *J. Satisf. Boolean Model. Comput.* 7 (2010), 59–6.
- [11] D.L. Bowen, L.H. Byrd, and William Clocksin. 1983. A portable Prolog compiler. In *D.A.I. research paper*. 9.
- [12] K. Dewey, J. Roesch, and B. Hardekopf. 2015. Fuzzing the Rust Type-checker Using CLP (T). In *Proceedings of the International Conference on Automated Software Engineering*. IEEE Computer Society, 482–493.
- [13] Yu Feng, Ruben Martins, Yuepeng Wang, Isil Dillig, and Thomas W. Reps. 2017. Component-based synthesis for complex APIs. In *Proceedings of the ACM SIGPLAN Symposium on Principles of Programming Languages*. ACM, 599–612.
- [14] Zheng Guo, Michael James, David Justo, Jiaxiao Zhou, Ziteng Wang, Ranjit Jhala, and Nadia Polikarpova. 2019. Program Synthesis by Type-Guided Abstraction Refinement. *Proc. ACM Program. Lang.* 4, POPL (2019), 12:1–12:28.
- [15] HyungSeok Han, DongHyeon Oh, and Sang Kil Cha. 2019. CodeAlchemist: Semantics-Aware Code Generation to Find Vulnerabilities in JavaScript Engines. In *Proceedings of the Network and Distributed System Security Symposium*. The Internet Society.
- [16] Christian Holler, Kim Herzig, and Andreas Zeller. 2012. Fuzzing with Code Fragments. In *Proceedings of the USENIX Security Symposium*. USENIX Association, USA, 445–458.
- [17] Kyriakos Ispoglou, Daniel Austin, Vishwath Mohan, and Mathias Payer. 2020. FuzzGen: Automatic Fuzzer Generation. In *Proceedings of the USENIX Security Symposium*. USENIX Association, 2271–2287.
- [18] Susmit Jha, Sumit Gulwani, Sanjit A. Seshia, and Ashish Tiwari. 2010. Oracle-Guided Component-Based Program Synthesis. In *Proceedings of the International Conference on Software Engineering*. Association for Computing Machinery, 215–224.
- [19] Ralf Jung, Hoang-Hai Dang, Jeehoon Kang, and Derek Dreyer. 2019. Stacked Borrows: An Aliasing Model for Rust. *Proc. ACM Program. Lang.* 4, POPL (2019), 41:1–41:32.
- [20] Y. Kim, Y. Kim, Taeksu Kim, Gunwoo Lee, Y. Jang, and M. Kim. 2013. Automated unit testing of large industrial embedded software using concolic testing. In *Proceedings of the International Conference on Automated Software Engineering*. IEEE, 519–528.
- [21] L. Ma, C. Artho, C. Zhang, H. Sato, J. Gmeiner, and R. Ramler. 2015. GRT: An Automated Test Generator Using Orchestrated Program Analysis. In *Proceedings of the International Conference on Automated Software Engineering*. IEEE Computer Society, 842–847.
- [22] Scott Olson. [n.d.]. *Miri*. <https://github.com/rust-lang/miri>
- [23] Carlos Pacheco, Shuvendu Lahiri, Michael D. Ernst, and Thomas Ball. 2006. *Feedback-directed Random Test Generation*. Technical Report MSR-TR-2006-125. Massachusetts Institute of Technology. 14 pages. <https://www.microsoft.com/en-us/research/publication/feedback-directed-random-test-generation/>
- [24] Boqin Qin, Yilun Chen, Zeming Yu, Linhai Song, and Yiyang Zhang. 2020. Understanding memory and thread safety practices and issues in real-world Rust programs. In *Proceedings of the ACM SIGPLAN International Conference on Programming Language Design and Implementation*. ACM, 763–779.
- [25] Eric Reed. 2015. *Patina: A formalization of the Rust programming language*. Technical Report. University of Washington.
- [26] Takehide Soh, Daniel Le Berre, Stéphanie Roussel, Mutsunori Banbara, and Naoyuki Tamura. 2014. Incremental SAT-Based Method with Native Boolean Cardinality Handling for the Hamiltonian Cycle Problem. In *Proceedings of the European Conference on Logics in Artificial Intelligence (Lecture Notes in Computer Science, Vol. 8761)*. Springer, 684–693.
- [27] John Toman, Stuart Pernsteiner, and Emina Torlak. 2015. CRUST: A Bounded Verifier for Rust. In *Proceedings of the International Conference on Automated Software Engineering*. IEEE Press, 6 pages.
- [28] F. Wang, F. Song, M. Zhang, X. Zhu, and J. Zhang. 2018. KRust: A Formal Executable Semantics of Rust. In *Proceedings of the International Symposium on Theoretical Aspects of Software Engineering*. IEEE Computer Society, 44–51.
- [29] Aaron Weiss, Daniel Patterson, Nicholas D. Matsakis, and Amal Ahmed. 2019. Oxide: The Essence of Rust. *CoRR* abs/1903.00982 (2019). arXiv:1903.00982 <http://arxiv.org/abs/1903.00982>
- [30] Maverick Woo, Sang Kil Cha, Samantha Gottlieb, and David Brumley. 2013. Scheduling Black-Box Mutational Fuzzing. In *Proceedings of the ACM SIGSAC Conference on Computer & Communications Security*. Association for Computing Machinery, 511–522.
- [31] Hui Xu, Zhuangbin Chen, Mingshen Sun, and Yangfan Zhou. 2020. Memory-Safety Challenge Considered Solved? An Empirical Study with All Rust CVEs. arXiv:2003.03296 [cs.PL]
- [32] Xuejun Yang, Yang Chen, Eric Eide, and John Regehr. 2011. Finding and understanding bugs in C compilers. In *Proceedings of the ACM SIGPLAN International Conference on Programming Language Design and Implementation*. ACM, 283–294.
- [33] Andreas Zeller, Rahul Gopinath, Marcel Böhme, Gordon Fraser, and Christian Holler. 2019. Fuzzing APIs. In *The Fuzzing Book*. Saarland University. <https://www.fuzzingbook.org/html/APIFuzzer.html> Retrieved 2019-12-17 16:45:28+01:00.